Document 4: Transformer Models and Self-Attention

1. The Transformer Architecture
Introduced in the paper "Attention Is All You Need", Transformers replaced recurrent networks by relying entirely on self-attention mechanisms to compute representations of input tokens.

2. Multi-Head Attention
Allows the model to jointly attend to information from different representation subspaces at different positions, capturing rich contextual dependencies across long text sequences.